

video:

https://drive.google.com/file/d/1F5oyxLjeRN3LDYVZyohvyjIVkG_GGkDv/view?usp=sharing



Aim

This project aims to build a machine learning model to predict the price of Airbnb listings based on various features such as property type, room type, location, amenities, and host characteristics. By analyzing these factors, this project will provide actionable insights to Airbnb hosts to optimize their listing prices.

Project- Part A: Airbnb Price Prediction and Insights

Importing Modules

In [138...

```
# importing required modules
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

import warnings
warnings.filterwarnings("ignore")
```

Loading Data

In [140...


```
# Load the data set
df = pd.read_csv('airbnb_data.csv')

# to display first five rows
df.head()
```

Out[140...

	id	log_price	property_type	room_type	amenities	accommodates
0	6901257	5.010635	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitch...	3
1	6304928	5.129899	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitch...	7
2	7919400	4.976734	Apartment	Entire home/apt	{TV,"Cable TV","Wireless Internet","Air condit...	5
3	13418779	6.620073	House	Entire home/apt	{TV,"Cable TV",Internet,"Wireless Internet",Ki...	4
4	3808709	4.744932	Apartment	Entire home/apt	{TV,Internet,"Wireless Internet","Air conditio...	2

5 rows × 29 columns



Data Exploration and Preprocessing

Data Description

In [142...

```
# display basic information
df.info()

# summary of statistics of the data set
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74111 entries, 0 to 74110
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     74111 non-null  int64
1   log_price                             74111 non-null  float64
2   property_type                         74111 non-null  object
3   room_type                             74111 non-null  object
4   amenities                             74111 non-null  object
5   accommodates                           74111 non-null  int64
6   bathrooms                             73911 non-null  float64
7   bed_type                              74111 non-null  object
8   cancellation_policy                   74111 non-null  object
9   cleaning_fee                          74111 non-null  bool
10  city                                   74111 non-null  object
11  description                            74111 non-null  object
12  first_review                           58247 non-null  object
13  host_has_profile_pic                   73923 non-null  object
14  host_identity_verified                 73923 non-null  object
15  host_response_rate                     55812 non-null  object
16  host_since                             73923 non-null  object
17  instant_bookable                       74111 non-null  object
18  last_review                           58284 non-null  object
19  latitude                               74111 non-null  float64
20  longitude                              74111 non-null  float64
21  name                                   74111 non-null  object
22  neighbourhood                           67239 non-null  object
23  number_of_reviews                      74111 non-null  int64
24  review_scores_rating                   57389 non-null  float64
25  thumbnail_url                          65895 non-null  object
26  zipcode                                73143 non-null  object
27  bedrooms                               74020 non-null  float64
28  beds                                   73980 non-null  float64
dtypes: bool(1), float64(7), int64(3), object(18)
memory usage: 15.9+ MB
```

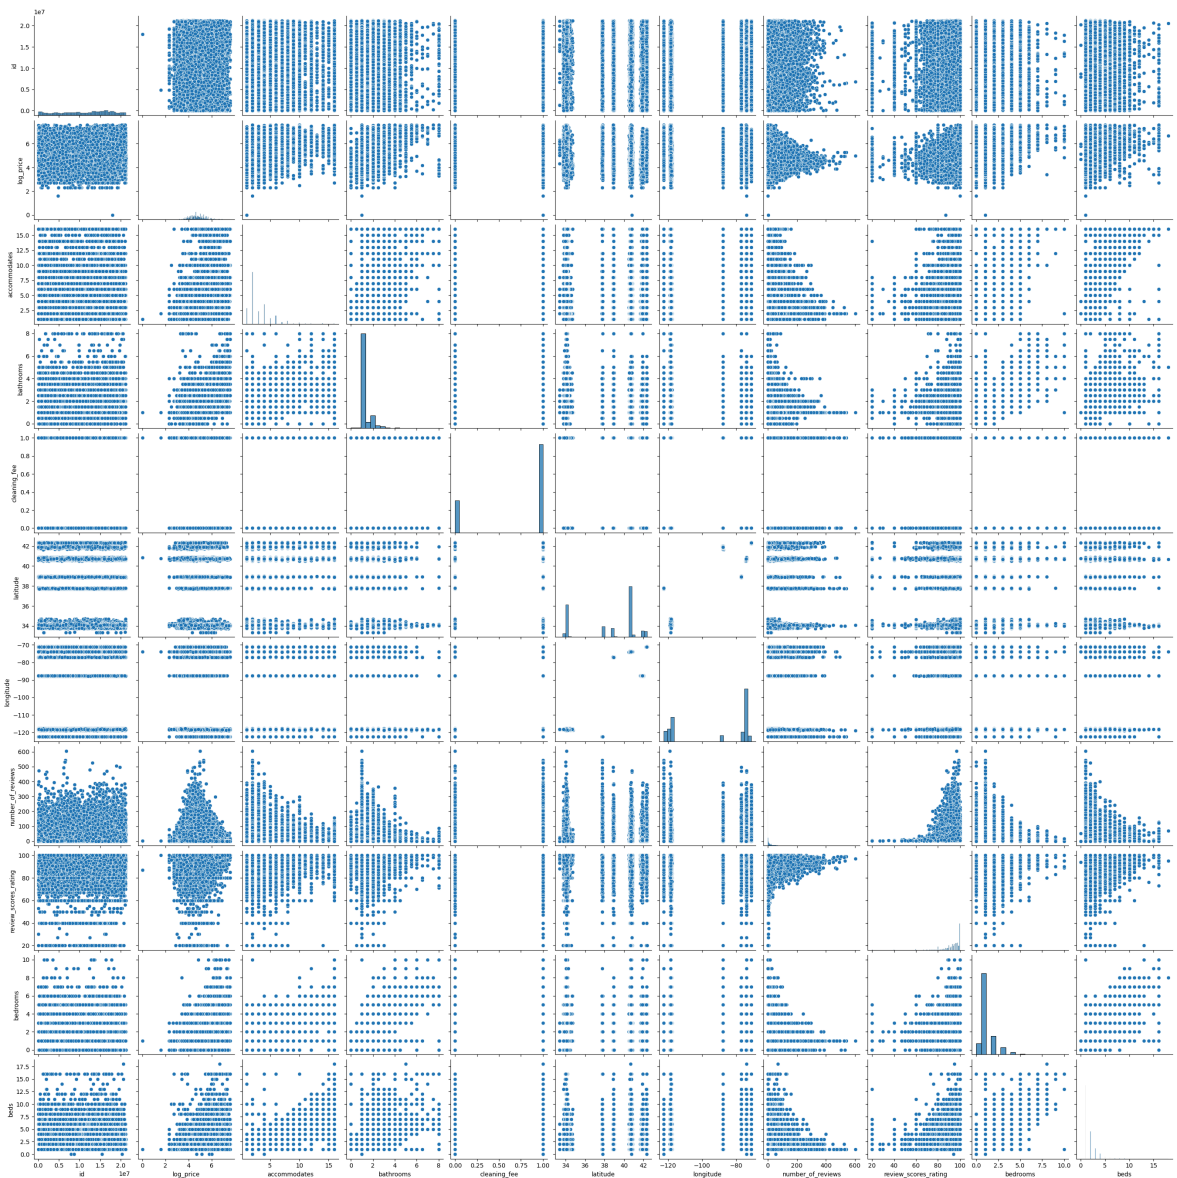
Out[142...

	id	log_price	accommodates	bathrooms	latitude	long
count	7.411100e+04	74111.000000	74111.000000	73911.000000	74111.000000	74111.0
mean	1.126662e+07	4.782069	3.155146	1.235263	38.445958	-92.3
std	6.081735e+06	0.717394	2.153589	0.582044	3.080167	21.7
min	3.440000e+02	0.000000	1.000000	0.000000	33.338905	-122.5
25%	6.261964e+06	4.317488	2.000000	1.000000	34.127908	-118.3
50%	1.225415e+07	4.709530	2.000000	1.000000	40.662138	-76.9
75%	1.640226e+07	5.220356	4.000000	1.000000	40.746096	-73.9
max	2.123090e+07	7.600402	16.000000	8.000000	42.390437	-70.9

Analyze the dataset for trends

In [144...


```
#pair plot visualizes relationships between variables.
sns.pairplot(df)#data
plt.show()
```



In [145...

```
# correlation matrix
corr_matrix = df[['id', 'log_price', 'accommodates', 'bathrooms', 'latitude', 'longit',
                  'review_scores_rating', 'bedrooms', 'beds']].corr()

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Matrix")
plt.show()
```



Missing Values

```
In [147... #missing values in columns
df.isnull().sum()
```

```
Out[147... id 0
log_price 0
property_type 0
room_type 0
amenities 0
accommodates 0
bathrooms 200
bed_type 0
cancellation_policy 0
cleaning_fee 0
city 0
description 0
first_review 15864
host_has_profile_pic 188
host_identity_verified 188
host_response_rate 18299
host_since 188
instant_bookable 0
last_review 15827
latitude 0
longitude 0
name 0
neighbourhood 6872
number_of_reviews 0
review_scores_rating 16722
thumbnail_url 8216
zipcode 968
bedrooms 91
beds 131
dtype: int64
```

Data cleaning

```
In [149... # columns left with missing values float64(contineous or Numeric). so, replaced
columns_to_fill = df[['bathrooms','review_scores_rating', 'bedrooms', 'beds']]

# Replace NaN values with mean
for column in columns_to_fill:
    mean = df[column].mean() # Find mean value
    df[column].replace(np.NaN, mean, inplace=True)
    df[column]=df[column].astype('int64')
```

```
In [150... #Columns which are Object(categorical) going to replace 'NaN' with most 'Frequen

columns_to_fill = df[['first_review','host_has_profile_pic', 'host_identity_veri
                    'last_review', 'neighbourhood', 'thumbnail_url',
                    'zipcode']]

# Replace NaN values with the most freq value
for column in columns_to_fill:
    freq = df[column].value_counts().idxmax() # Find the most frequent value
    df[column].replace(np.NaN, freq, inplace=True)
```

```
In [151... # convert Boolean (True or False) to binary form
df['cleaning_fee']=df['cleaning_fee'].astype('int64')

#By using map()
convert_columns = df[['host_has_profile_pic', 'host_identity_verified','instant_
```

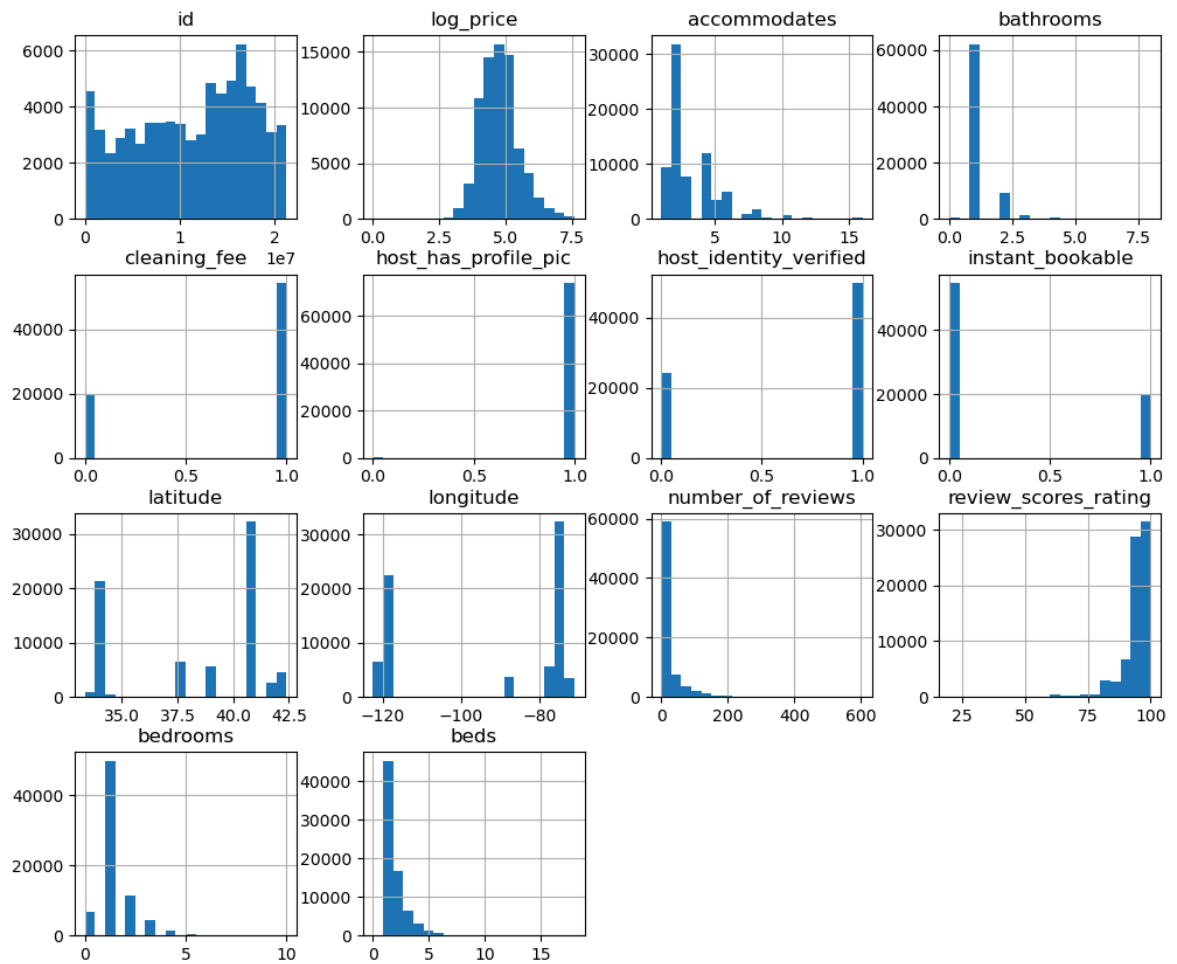
```
for column in convert_columns:
    df[column] = df[column].map({'t': 1, 'f': 0})
```

```
In [152... #checking missing values
df.isnull().sum()
```

```
Out[152... id                0
log_price             0
property_type         0
room_type             0
amenities             0
accommodates          0
bathrooms             0
bed_type              0
cancellation_policy   0
cleaning_fee          0
city                  0
description            0
first_review          0
host_has_profile_pic   0
host_identity_verified 0
host_response_rate     0
host_since            0
instant_bookable       0
last_review           0
latitude              0
longitude             0
name                  0
neighbourhood          0
number_of_reviews      0
review_scores_rating   0
thumbnail_url         0
zipcode               0
bedrooms              0
beds                  0
dtype: int64
```

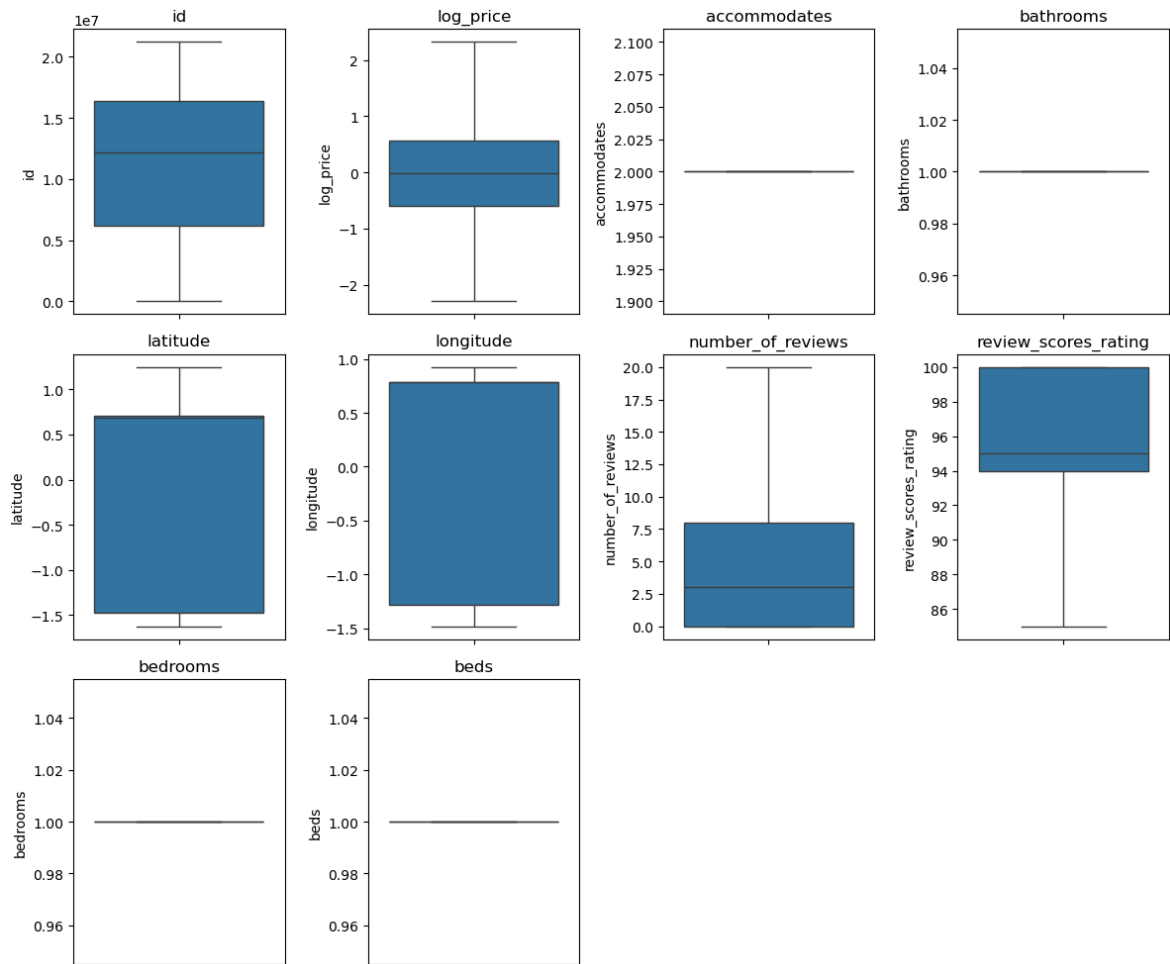
Outliers

```
In [154... # histogram for features.
df.hist(figsize=(12,10),bins=20)
plt.show()
```



```
In [216... # Box plot for outliers.
plt.figure(figsize=(12, 10))
for i, column in enumerate(df[['id', 'log_price', 'accommodates', 'bathrooms', 'latit',
                              'review_scores_rating', 'bedrooms', 'beds']]):

    plt.subplot(3,4,i+1)
    sns.boxplot(df[column])
    plt.title(column)
plt.tight_layout()
plt.show()
```

Remove outliers using IQR method

```
In [214... # Remove outliers using IQR method
column = df[['id', 'log_price', 'accommodates', 'bathrooms', 'latitude', 'longitude',
              'review_scores_rating', 'bedrooms', 'beds']]

for col in column:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    df = df[~((df[col] < (Q1 - 1.5 * IQR)) | (df[col] > (Q3 + 1.5 * IQR)))]
```

Feature Engineering

```
In [159... #converting amenities to numeric by length
df['amenities'] = df['amenities'].apply(lambda x: len(x.split(',')))
```

Transformations.

```
In [161... #Encoding
df = pd.get_dummies(df, columns=['room_type'], dtype=int)
```

```
In [162... from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['host_response_rate'] = le.fit_transform(df['host_response_rate'])
```

```

In [163... le = LabelEncoder()
df['property_type'] = le.fit_transform(df['property_type'])

In [164... from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df[['log_price', 'latitude', 'longitude']] = scaler.fit_transform(df[['log_price

In [165... df.columns

Out[165... Index(['id', 'log_price', 'property_type', 'amenities', 'accommodates',
      'bathrooms', 'bed_type', 'cancellation_policy', 'cleaning_fee', 'city',
      'description', 'first_review', 'host_has_profile_pic',
      'host_identity_verified', 'host_response_rate', 'host_since',
      'instant_bookable', 'last_review', 'latitude', 'longitude', 'name',
      'neighbourhood', 'number_of_reviews', 'review_scores_rating',
      'thumbnail_url', 'zipcode', 'bedrooms', 'beds',
      'room_type_Entire home/apt', 'room_type_Private room',
      'room_type_Shared room'],
      dtype='object')

```

Model Development

Build a regression model to predict listing prices.

```

In [218... #define multiple predictor
x_multi = df[['property_type','amenities','accommodates','bathrooms','cleaning_f
      'host_response_rate','instant_bookable','latitude','longitude','number_of_re
      'bedrooms','beds','room_type_Entire home/apt','room_type_Private room','room
y_multi = df['log_price']

In [220... # Split the dataset into training and testing sets
x_train_multi,x_test_multi,y_train_multi,y_test_multi = train_test_split(x_multi
      test_siz

In [222... #create and traun the model
multi_model = LinearRegression()
multi_model.fit(x_train_multi,y_train_multi)

Out[222... ▼ LinearRegression ⓘ ?
LinearRegression()

In [224... #predication
y_pred_multi = multi_model.predict(x_test_multi)

#Evaluate the model
print(f"Mean Squared Error: {mean_squared_error(y_test_multi,y_pred_multi)}")
print(f"R^ Score: {r2_score(y_test_multi,y_pred_multi)}")

Mean Squared Error: 0.5339284962033279
R^ Score: 0.31575991152572835

In [226... # Display the coefficients of the model
coefficients = pd.DataFrame(multi_model.coef_.flatten(), x_multi.columns, column

```

```
print(coefficients)
```

	Coefficient
property_type	-1.263407e-03
amenities	9.978617e-03
accommodates	1.110223e-16
bathrooms	-3.885781e-16
cleaning_fee	-2.234336e-02
host_has_profile_pic	-4.732447e-02
host_identity_verified	2.380934e-03
host_response_rate	-6.856752e-04
instant_bookable	-1.160430e-01
latitude	2.748607e-01
longitude	-2.446982e-01
number_of_reviews	-6.448629e-03
review_scores_rating	4.702135e-03
bedrooms	0.000000e+00
beds	0.000000e+00
room_type_Entire home/apt	7.586729e-01
room_type_Private room	-1.895214e-01
room_type_Shared room	-5.691515e-01
number_of_reviews	-6.448629e-03


In [228...

```
# Predicted vs. Actual Plot
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test_multi, y=y_pred_multi, alpha=0.6, color='blue')
plt.plot([y_test_multi.min(), y_test_multi.max()], [y_test_multi.min(), y_test_multi.max()],
         color='red', linestyle='--', linewidth=2)
plt.title('Predicted vs Actual Prices')
plt.xlabel('Actual Prices (log_price)')
plt.ylabel('Predicted Prices (log_price)')
plt.show()


# Residuals Plot
residuals = y_test_multi - y_pred_multi
plt.figure(figsize=(8, 6))
sns.histplot(residuals, kde=True, color='purple')
plt.title('Residuals Distribution')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.show()


# Residuals vs. Predicted Plot
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_pred_multi, y=residuals, alpha=0.6, color='green')
plt.axhline(0, color='red', linestyle='--', linewidth=2)
plt.title('Residuals vs Predicted Prices')
plt.xlabel('Predicted Prices (log_price)')
plt.ylabel('Residuals')
plt.show()
```

Predicted vs Actual Prices



Residuals Distribution





In []:

In []: